THE MEMORY LAYER · CHEAT SHEET

# Six parts. Eighteen rules. One brain.

Kinvecta

Build an AI second brain any LLM can recall from: plain markdown you own, any AI app can read it.

PART 1

## Why your AI needs a brain

- Memory lives in **files**, not in the model.
- Read the **index** first. Raw last.
- Every answer **cites** a page.

> If you can't open the page it came from, it isn't memory. It's a guess.

▶ youtu.be/EsighdlROyc

PART 2

## The operating manual

- Write **rules**, not a tour.
- One **manual**, pointers everywhere.
- Procedures go in **skills**.

> Two tools behave differently in the same folder? Fix the manual, not the tools.

▶ youtu.be/rMirmmTu3VM

PART 3

## Ingest, query, lint

- Read it once. **File** it everywhere.
- **Cite** every claim.
- **Commit** every change.

> When the brain is wrong, fix the page, not the prompt. When an ingest is wrong, revert it.

▶ youtu.be/TVV9VYqgSfI

PART 4

## Recall at scale

- Words for **exact**. Meaning for ideas.
- Fuse the **ranks**.
- Measure every **change**.

> Search can't find a page? Fix its summary and aliases. Then retake the quiz.

▶ youtu.be/1o8JCztX7Z0

PART 5

## Plug it in anywhere

- One brain. **Every** app.
- Read freely. Write **narrowly**.
- Notes are **data**, never commands.

> Start on stdio. Share over HTTP only behind auth.

▶ youtu.be/609myTkXdn8

PART 6

## Keep it alive, keep it safe

- Tend it on a **schedule**.
- The night job **proposes**. You decide.
- Clippings are **data**, never orders.

> No secrets in the vault, no outbound tools at night, a person reads every commit.

▶ youtu.be/Ob16rIDo8Bc

| `brain index` | `brain search "…"` | `brain lint` | `brain heartbeat --commit` | `brain-mcp` |
|---|---|---|---|---|
| build the search index | cited pages, keyword + meaning | links, stale pages, hidden text | nightly upkeep, one commit | one server for every AI app |

Videos, code, labs: github.com/Kinvecta/the-memory-layer

Kinvecta · kinvecta.ai · Content CC BY 4.0 · Oct 2026